

An aerial photograph of a suburban neighborhood, likely in a tropical or subtropical region, characterized by numerous houses with dark tiled roofs and light-colored walls. Palm trees and other lush greenery are interspersed among the buildings. The scene is captured from a high angle, showing the layout of the streets and the density of the housing.

NYCDSA Machine Learning Project: Ames Housing Dataset

The All-American Regex



Agenda

- I. Introduction & Background
- II. Data Pre-Processing
- III. Linear Models
- IV. Tree-Based Models
- V. Stacked Models
- VI. Conclusions & Next Steps



Background

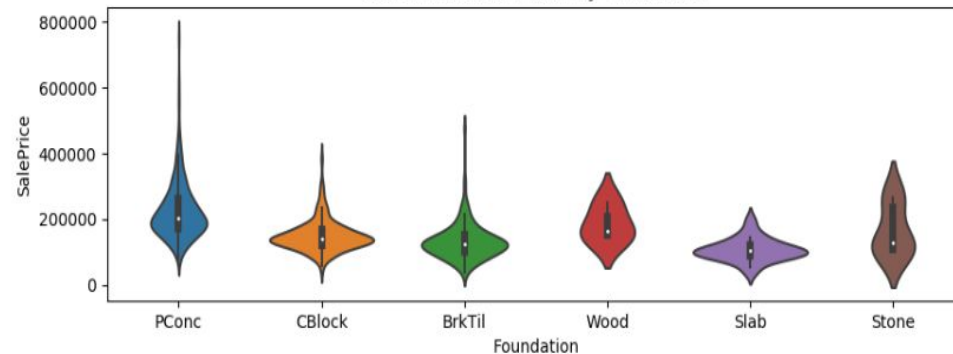
- ❑ Housing sales in Ames, Iowa from 2006 to 2010
- ❑ 2,919 sales (1,460 in the training set)
- ❑ 80 features (23 nominal, 23 ordinal, 14 discrete, 20 continuous)
 - ❑ Size, quality, area, year, etc.
- ❑ Originally deployed by Dean De Cock in 2011 as an alternative to the Boston Housing Dataset (Harrison and Rubenfeld 1978)

Features & Some EDA

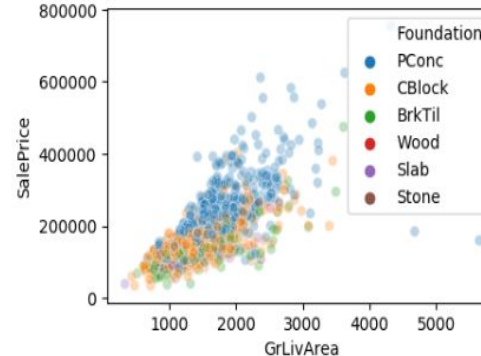
- Overall material and finish (OverallQual) and above ground square footage (GrLivArea) had the strongest linear relationships with the sale price
- We took a look at each feature with respect to the two above and brainstormed how we could feed them into a model

Choose var: Foundation

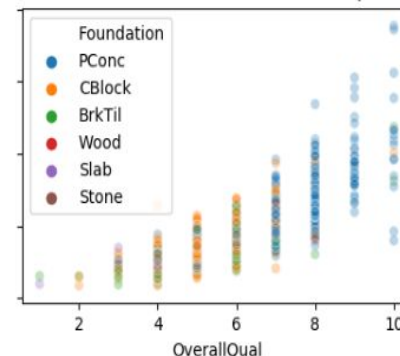
Sale Price Distribution by Foundation



Effect of Foundation after GrLivArea



Effect of Foundation after OverallQual



Pre-Processing: Exclusion and Imputation

- ❑ Some features appeared to be collinear with others or insignificant, so we decided to remove them from our data (eg. Utilities)
 - ❑ We ended up moving forward with 53 features
- ❑ Dealing with “NaN” values - Fill with “None”, 0, or the mode of the training set based on variable type
 - ❑ LotFrontage was filled with the neighborhood median of the training set
- ❑ Two outliers in the training set (>4,000SF, <\$200k) were removed

Pre-Processing: Handling Features of Different Types

❑ One-hot encoding

- ❑ Negligible minority classes were dropped

Foundation_BrkTil	Foundation_CBlock	Foundation_PConc
0	0	1
0	1	0
0	0	1
1	0	0
0	0	1
0	0	0
0	0	1
0	1	0

❑ Binary variables

- ❑ eg. Number of fireplaces → Is there a fireplace?

Fireplaces	PavedDrive
0	1
1	1
1	1
1	1
1	1
1	1
0	1
1	1
1	1
1	1


Pre-Processing: Special Cases

☐ Transformations

☐ eg. LotArea \rightarrow $\log(\text{LotArea})$


☐ Grouping

☐ eg. YearBuilt before 1950 \rightarrow 1950



Simple Linear Regression

House Size vs. Sale Price

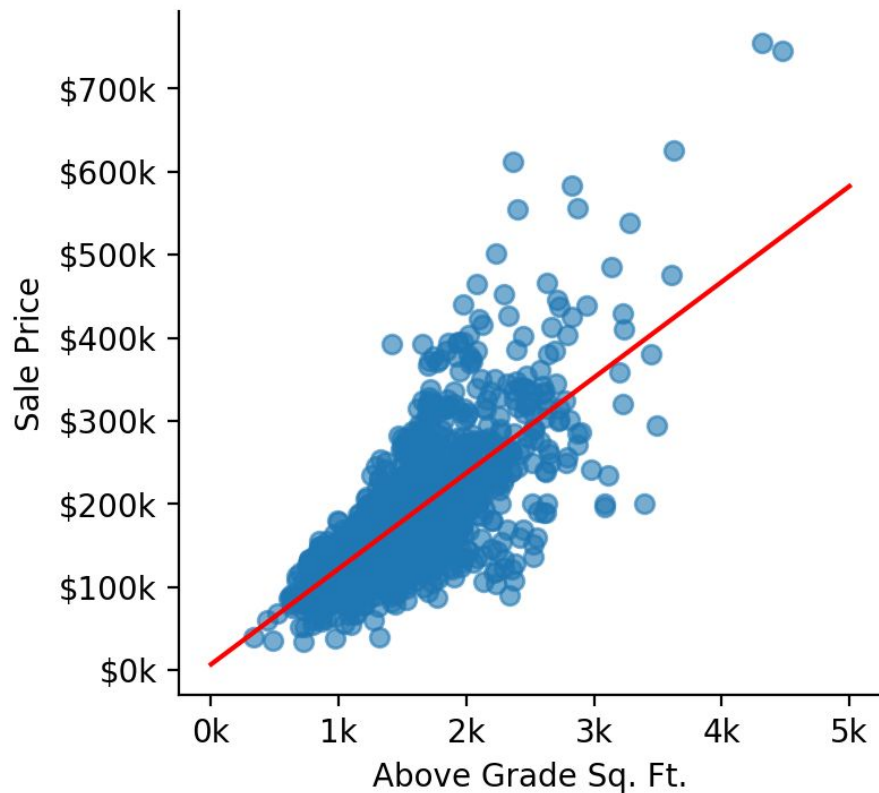


$$\text{SalePrice} = 7,165 + 115 * \text{GrLivArea}$$


- R^2 (Training) = 0.54
- RMSLE (Training) = 0.273
- RMSLE (CV) = 0.273

Simple Linear Regression: Residuals

House Size vs. Sale Price

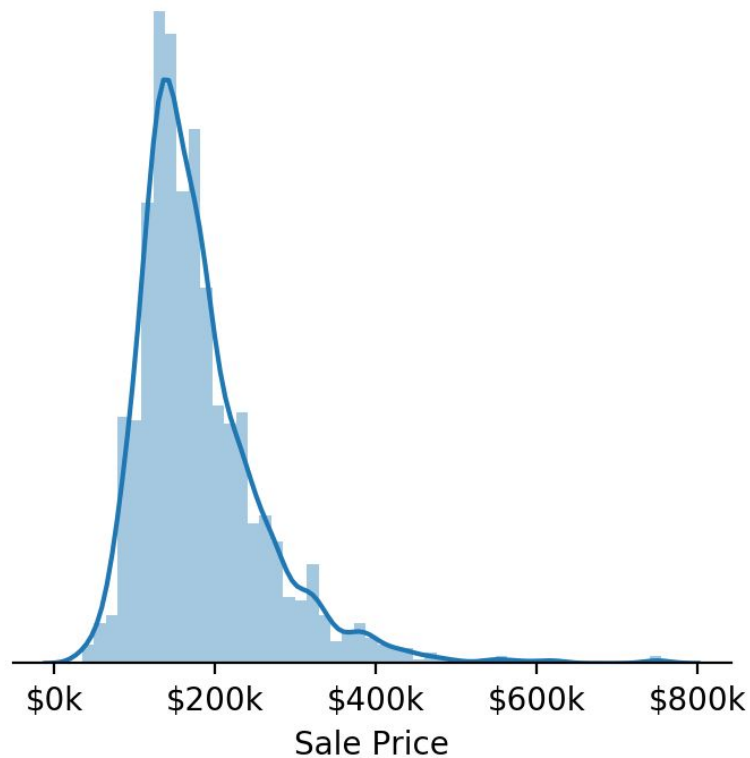


Residual Plot

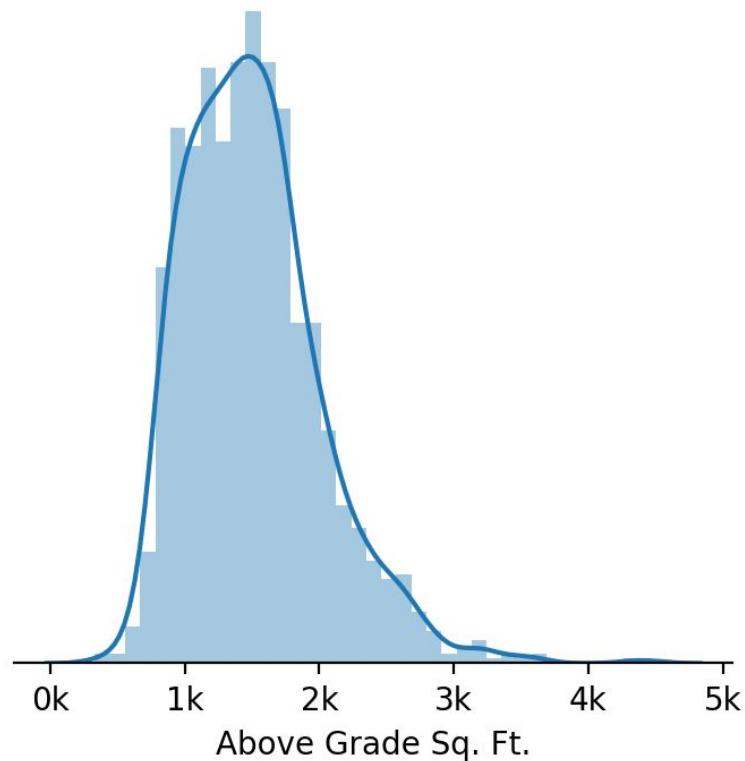


Simple Linear Regression: Variable Distributions

Distribution of Sale Price




Distribution of House Size



Simple Linear Regression: Box-Cox

House Size vs. Sale Price

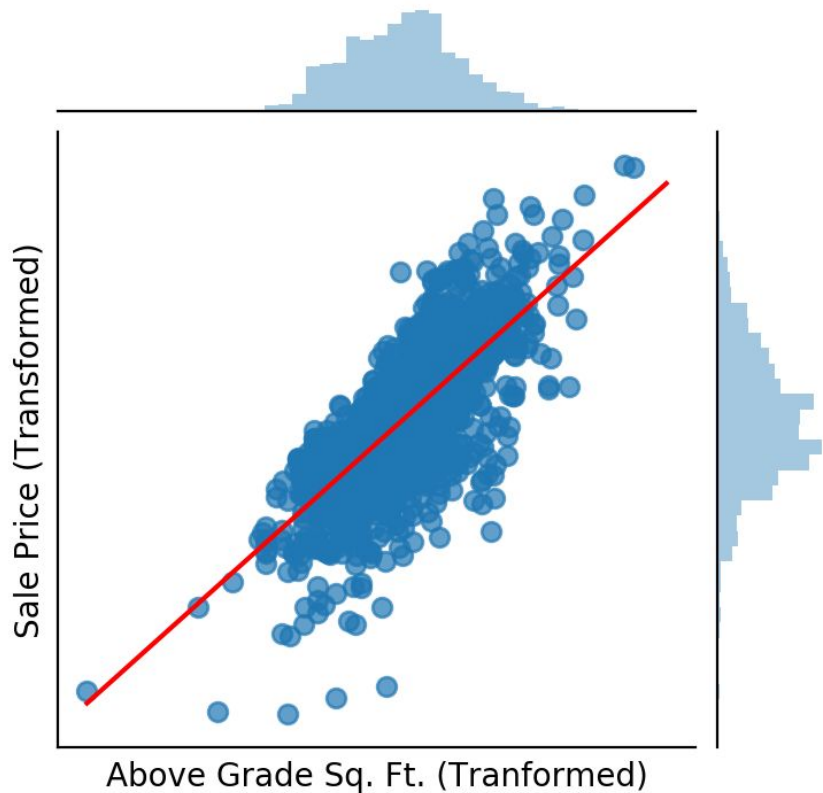


Transformed Model

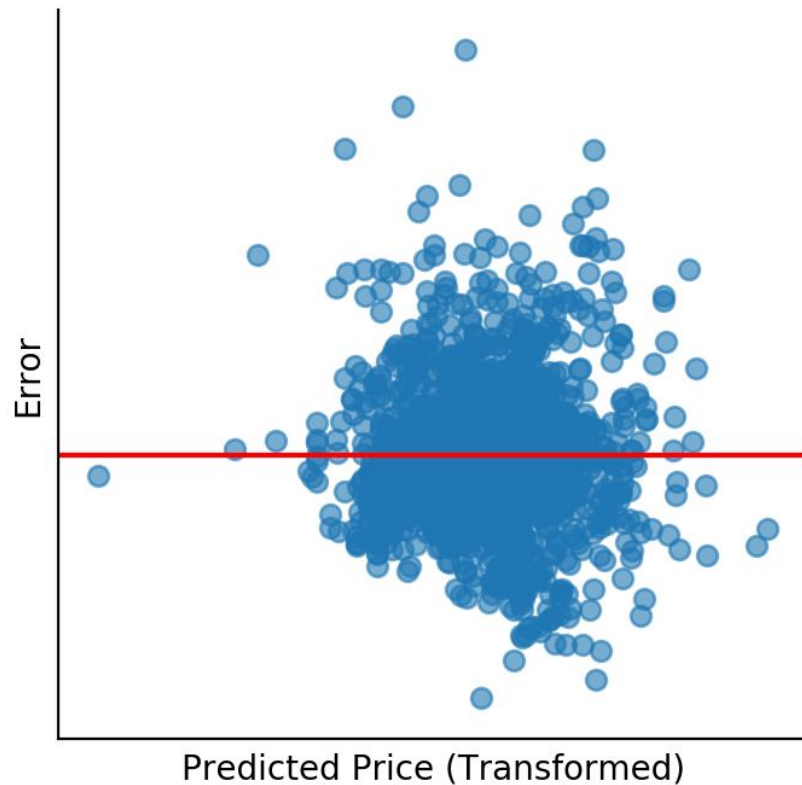
- R^2 (Training) = 0.54
- RMSLE (Training) = 0.269
- RMSLE (CV) = 0.270

Simple Linear Regression: Box-Cox

House Size vs. Sale Price



Residual Plot





Multiple Linear Regression



Multiple Linear Regression

L2 Regularization

L1 Regularization

No Regularization

- R^2 (Training) = 0.61
- RMSLE (Training) = 0.250
- RMSLE (CV) = 0.251

L2 Regularization (Ridge)

- R^2 (Training) = 0.92
- RMSLE (Training) = 0.108
- RMSLE (CV) = 0.119

L1 Regularization (Lasso)


- R^2 (Training) = 0.92
- RMSLE (Training) = 0.109
- RMSLE (CV) = 0.119

Combined (Elastic Net)

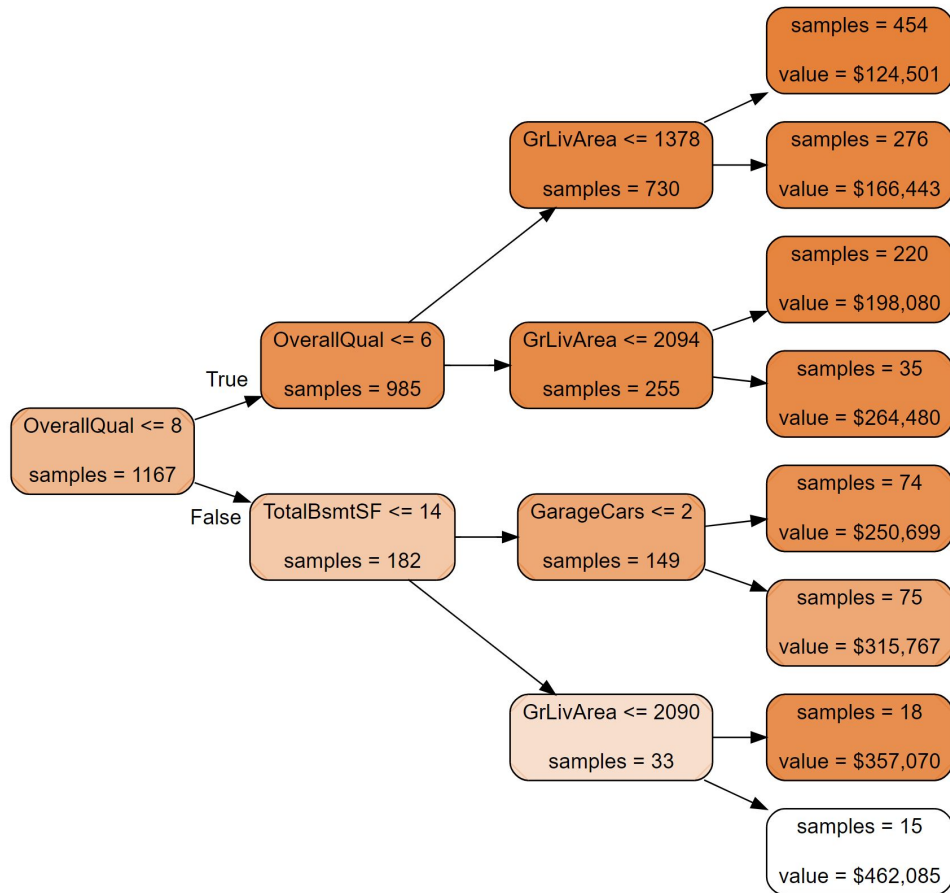
- R^2 (Training) = 0.92
- RMSLE (Training) = 0.111
- RMSLE (CV) = 0.118

Multiple Linear Regression

Elastic Net Variable Importance



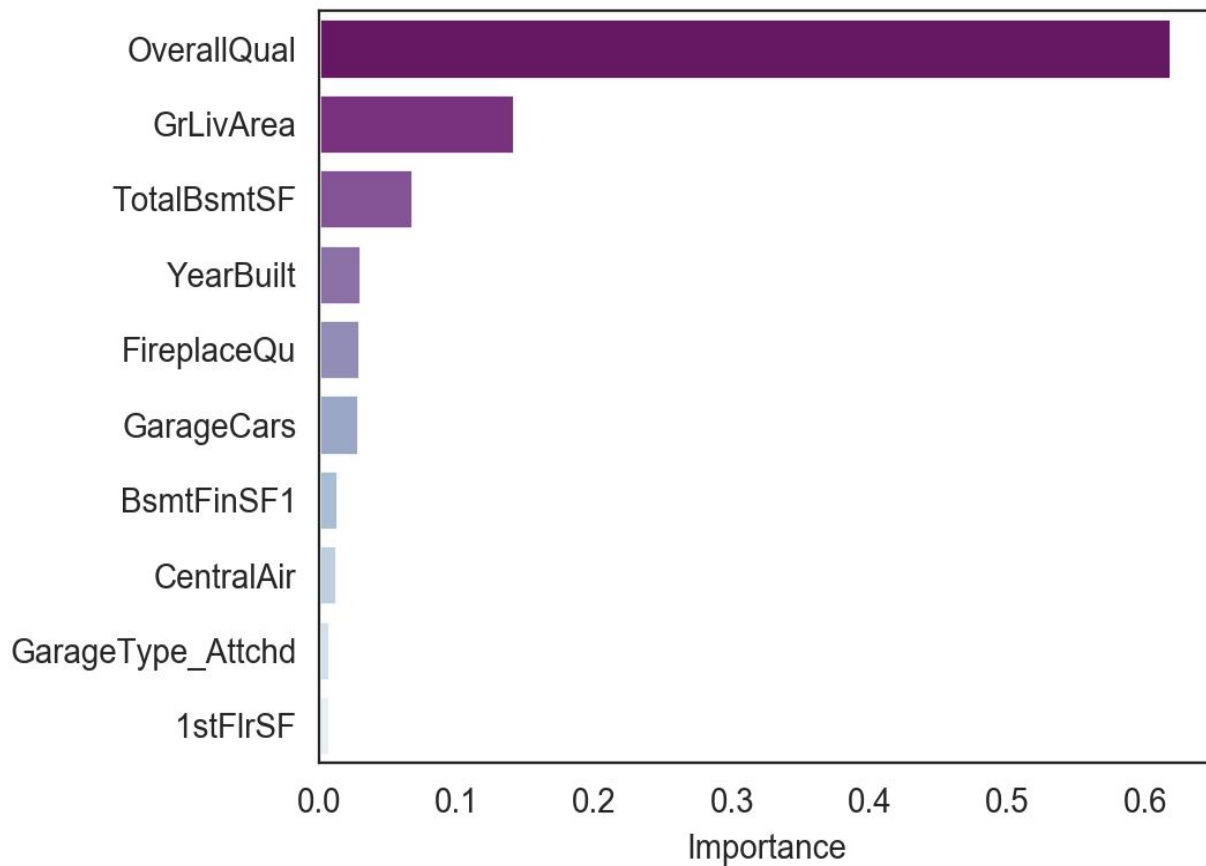
Tree-Based Models



Important Hyperparameters

- N_estimators
- Max_features
- Max_depth
- Min_samples_split
- Min_samples_leaf

Decision Tree

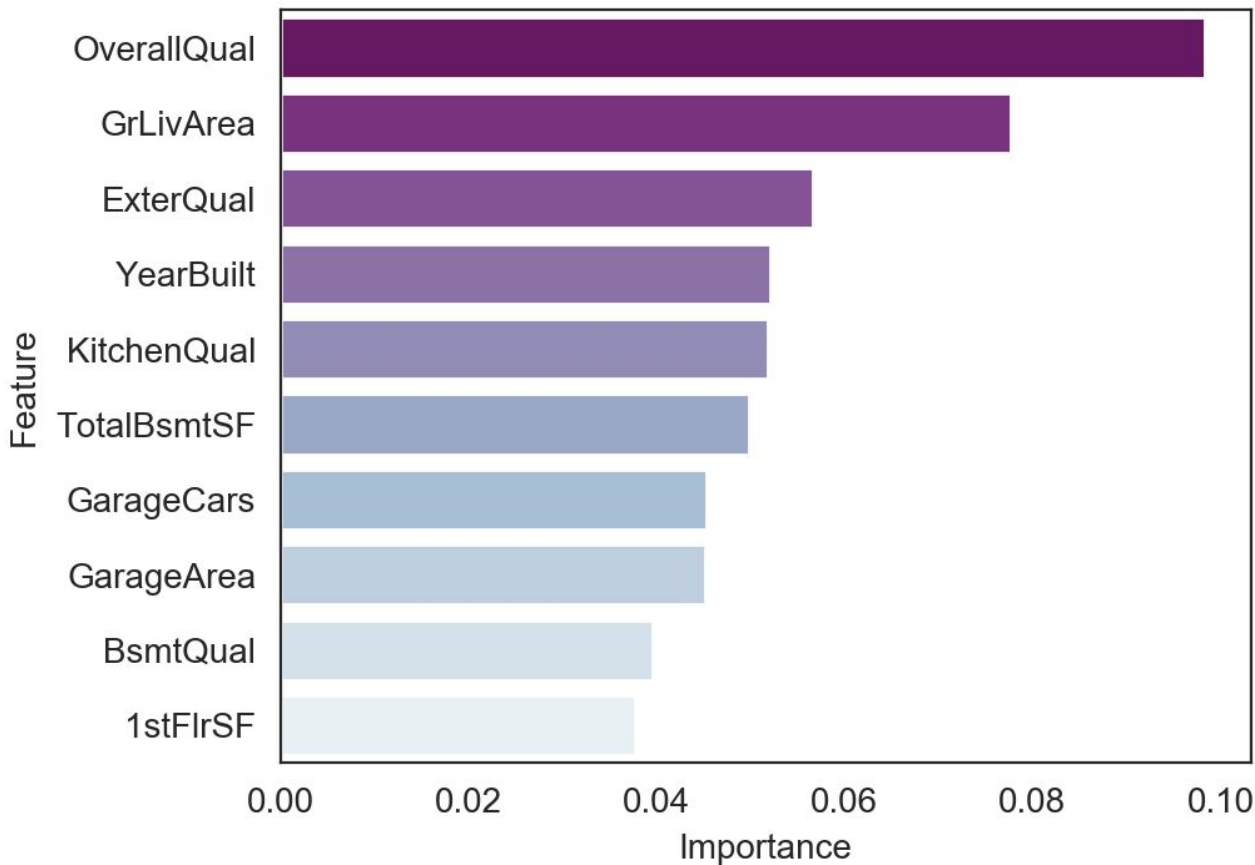


R^2 (Training) = 0.87

RMSLE (Training) = 0.141

RMSLE (CV) = 0.184

Random Forest



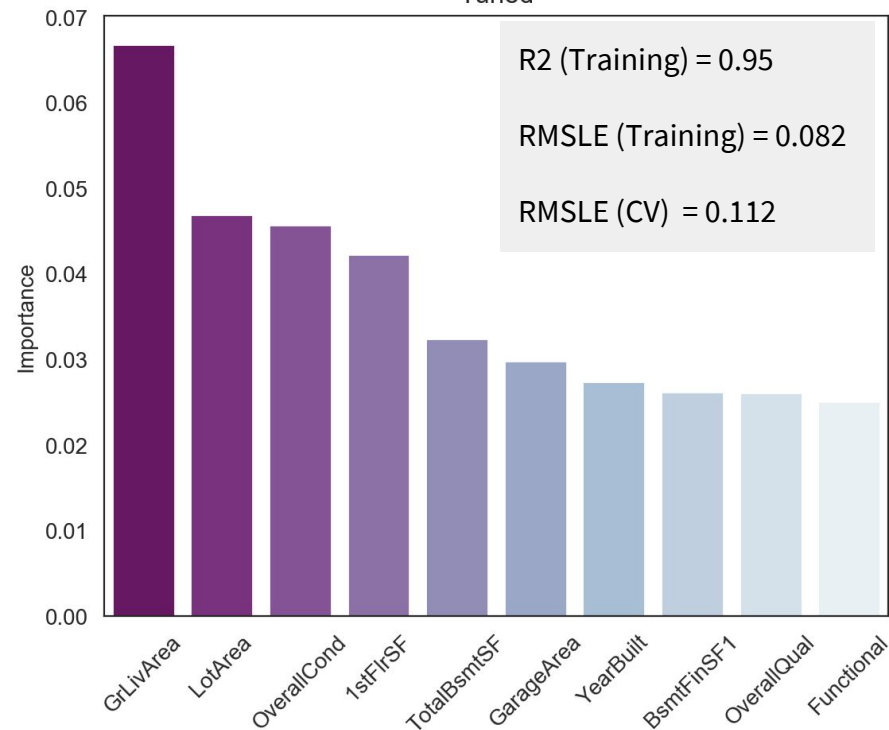
R^2 (Training) = 0.85

RMSLE (Training) = 0.151


RMSLE (CV) = 0.165

Gradient Boosting


Tuned



Untuned



Gradient Boosting

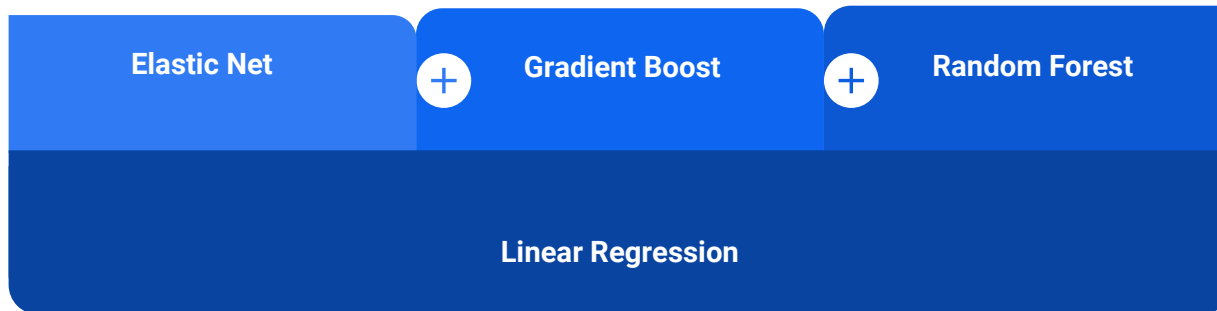


Exploring Ensembling

- Intuition
 - Elastic Net (PL) 0.121
 - Tuned Gradient Boost (PL) 0.122
 - Tuned Random Forest (PL) 0.145
- Averaging
 - Equal weight to top 3 performers (PL) 0.123
 - Drop weakest link (PL) **0.118**

Stacking

- Stacking
 - Elastic Net, Gradient Boost, Random Forest base learners
 - Linear regression meta-model (PL) **0.117**
 - Elastic Net, Gradient Boost
 - Linear regression meta-model (PL) **0.117**





Conclusions / Next Steps

- Importance of Feature Engineering
 - EDA → Feature Engineering
- Public vs. Private Leaderboard
- Future Work
 - Explore applicability and effectiveness of PCA and MCA
 - Ensembling
 - Stacking architecture (size of stack, strategic tuning parameters)
 - Different high-level learners